

Better Gradient Descent

I simulated some silly data to run the following algorithm.

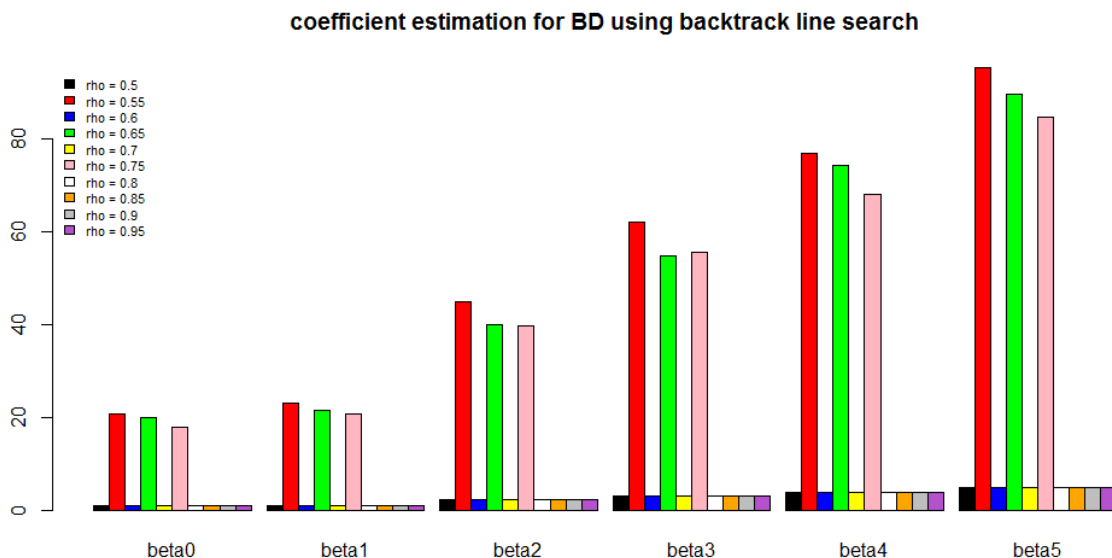
```
set.seed(128)
n = 1000
x1 = rnorm(n)
x2 = rnorm(n)
x3 = rnorm(n)
x4 = rnorm(n)
x5 = rnorm(n)
z = 1 + 1 * x1 + 2 * x2 + 3 * x3 + 4 * x4 + 5 * x5
pr = 1/(1+exp(-z))
y = rbinom(n,1,pr)
```


Backtrack line search

Performance with respect to rho (c = 0.001, alpha0 = 1)

The estimation does not work well for $\rho = 0.55$, $\rho = 0.65$, and $\rho = 0.75$. From the log likelihood value, the problem is that algorithms with those parameter values converge fast. Any idea why this is happening? I use log likelihood value as convergence criterion.

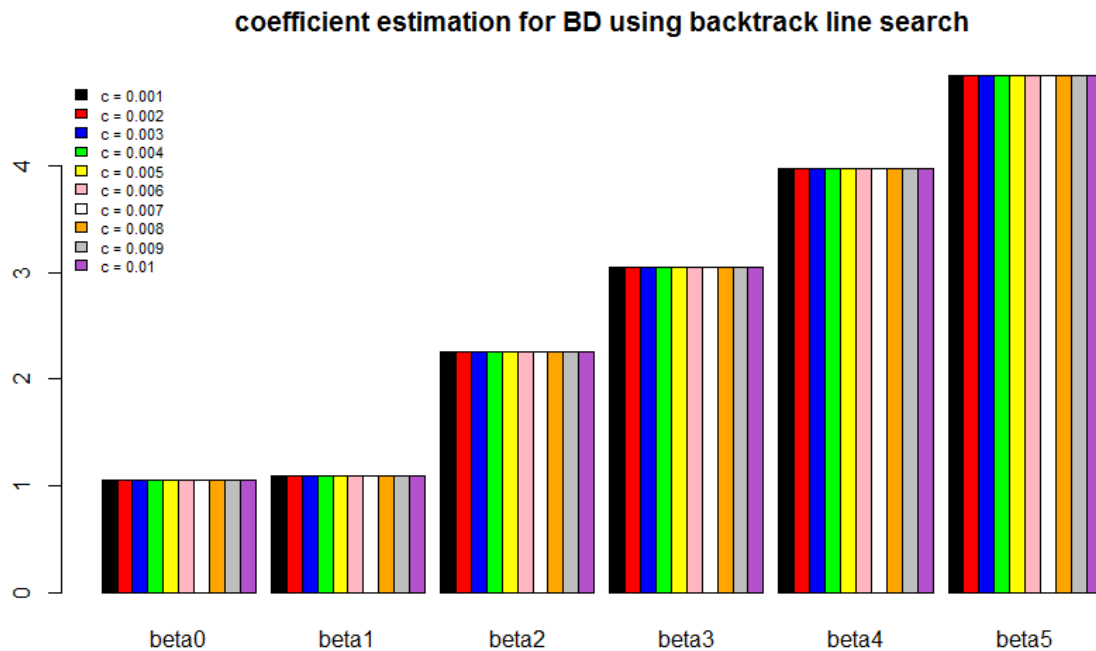
As ρ grows, it takes more iterations to converge. That makes sense since a big ρ leads to small change in step size, thus taking more steps to get there.

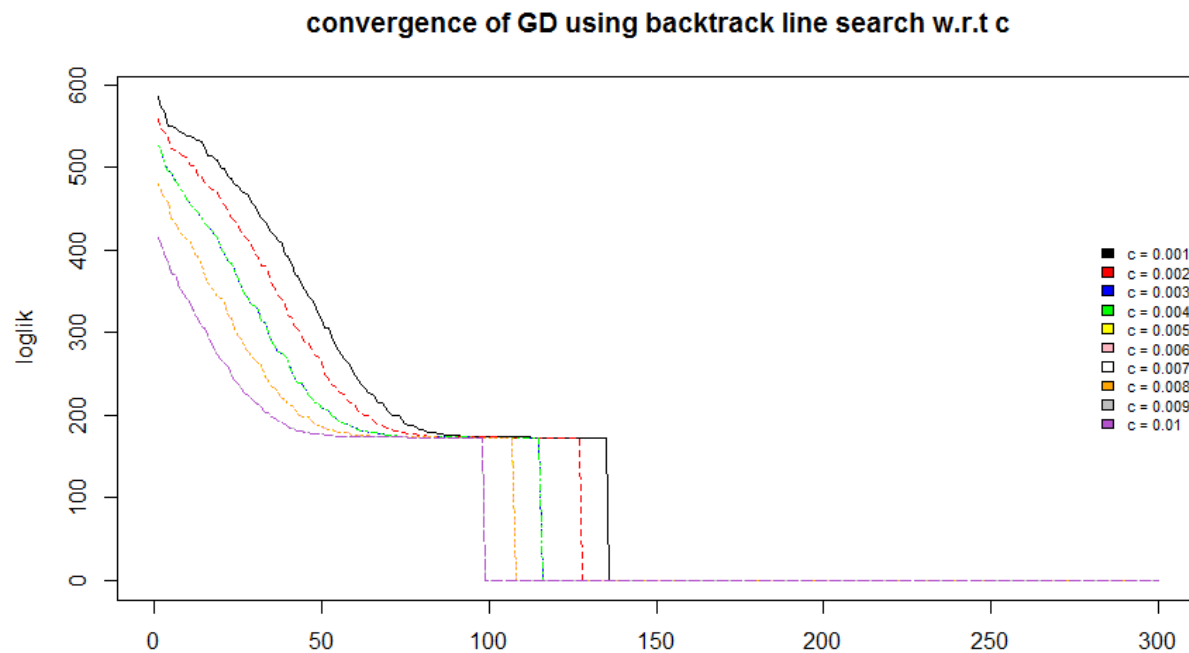




Performance with respect to c ($\rho = 0.7$, $\alpha_0 = 1$)

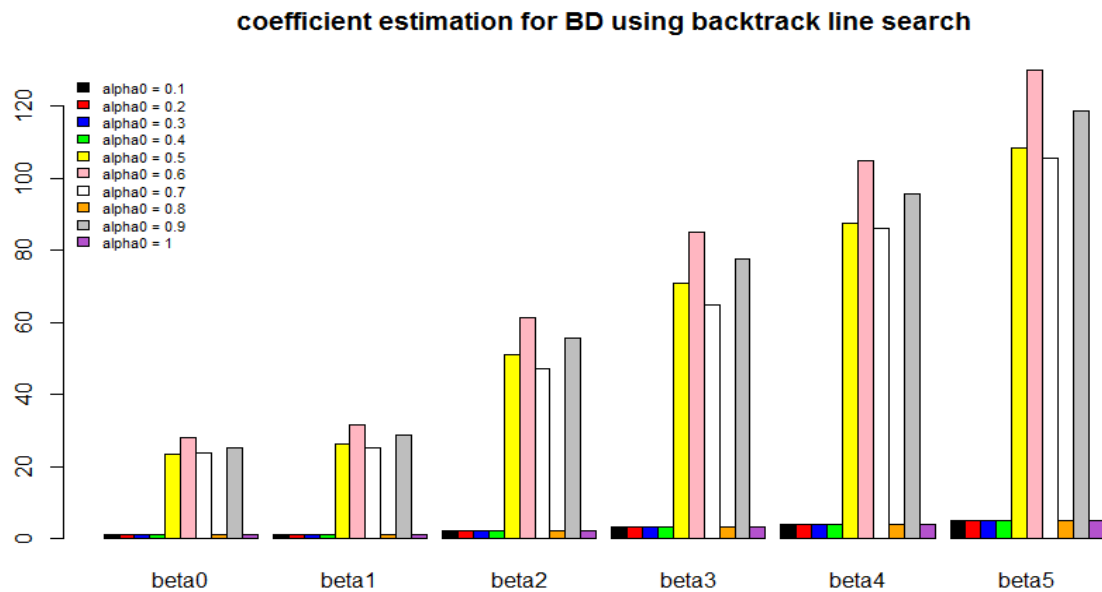
The estimation works well. As c grows, it takes fewer iterations to converge. Again, a small c leads to small change in step size in each iteration, thus taking more steps to get there.







Performance with respect to c ($\rho = 0.7$, $c = 0.001$)

The estimation breaks again for $\alpha_0 = 0.5, 0.6, 0.6, 0.9$. The problem is the same as before that it converges too fast. Any idea why this is happening? As α_0 grows, it takes more iterations to converge. Since it is backtrack line search, a big α_0 means we are initially far from the optimal α , thus taking more steps to get there.




convergence of GD using backtrack line search w.r.t initial alpha






Performance with respect to c ($\rho = 0.7$, $\alpha_0 = 1$)

coefficient estimation for Quasi-Newton using backtrack line search



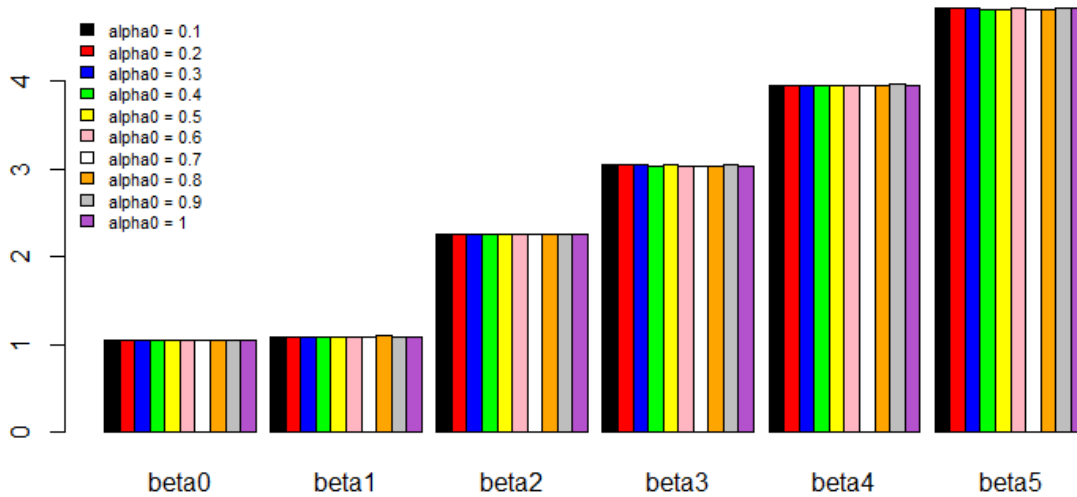
convergence of Quasi-Newton using backtrack line search w.r.t c




Performance with respect to initial alpha ($\rho = 0.7$, $c = 0.001$)

Again, the estimation problem seems fixed under Quasi Newton method.

coefficient estimation for Quasi-Newton using backtrack line search



convergence of Quasi-Newton using backtrack line search w.r.t initial alpha



GD v.s. GD using backtrack line search v.s. Quasi-Newton using backtrack line search

Parameter setting:

GD: $\alpha = 0.01$

GD using backtrack line search and Quasi-Newton using backtrack line search

$\text{Alpha0} = 1$, $c = 0.001$, $\rho = 0.7$

Test convergence based on log likelihood value. Threshold is set at $1e-5$.

Takeaway. All algorithms return satisfactory estimates. Incorporating backtrack line search into GD takes fewer iterations to converge. But, unexpectedly, Quasi-Newton takes more iterations to converge than GD, no matter it includes backtrack line search or not. The reason might be attributed to the selection of our parameters. In terms of the time each algorithm runs, GD is fastest with mean milliseconds of 70.32, and Quasi-Newton using backtrack line search is slowest with mean milliseconds of 1024.66. So, backtrack line search is time consuming, but gives better step size and converges in fewer iterations.

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
GD(X, y)	65.81369	67.54288	70.32293	68.33913	69.32054	122.2452	100	

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
qNewton(X, y)	359.8116	366.5836	377.04	369.327	379.7995	468.2446	100	

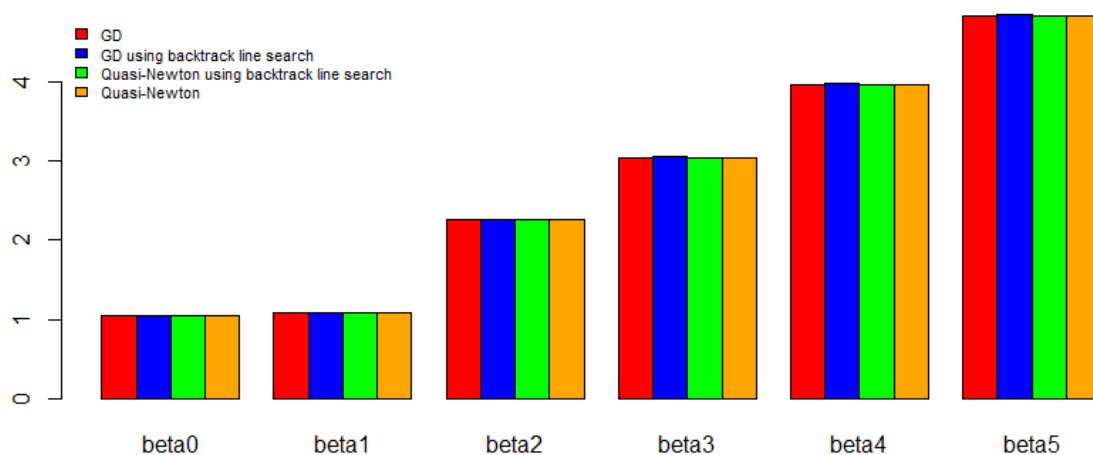
Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
BGD(X, y)	372.2502	375.6868	379.2561	377.5591	379.9752	428.5745	100	

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
qNewton_bb(X, y)	976.5969	984.299	1024.661	1025.465	1049.769	1122.962	100	

comparison of GD and its variants



convergence of GD and its variants

